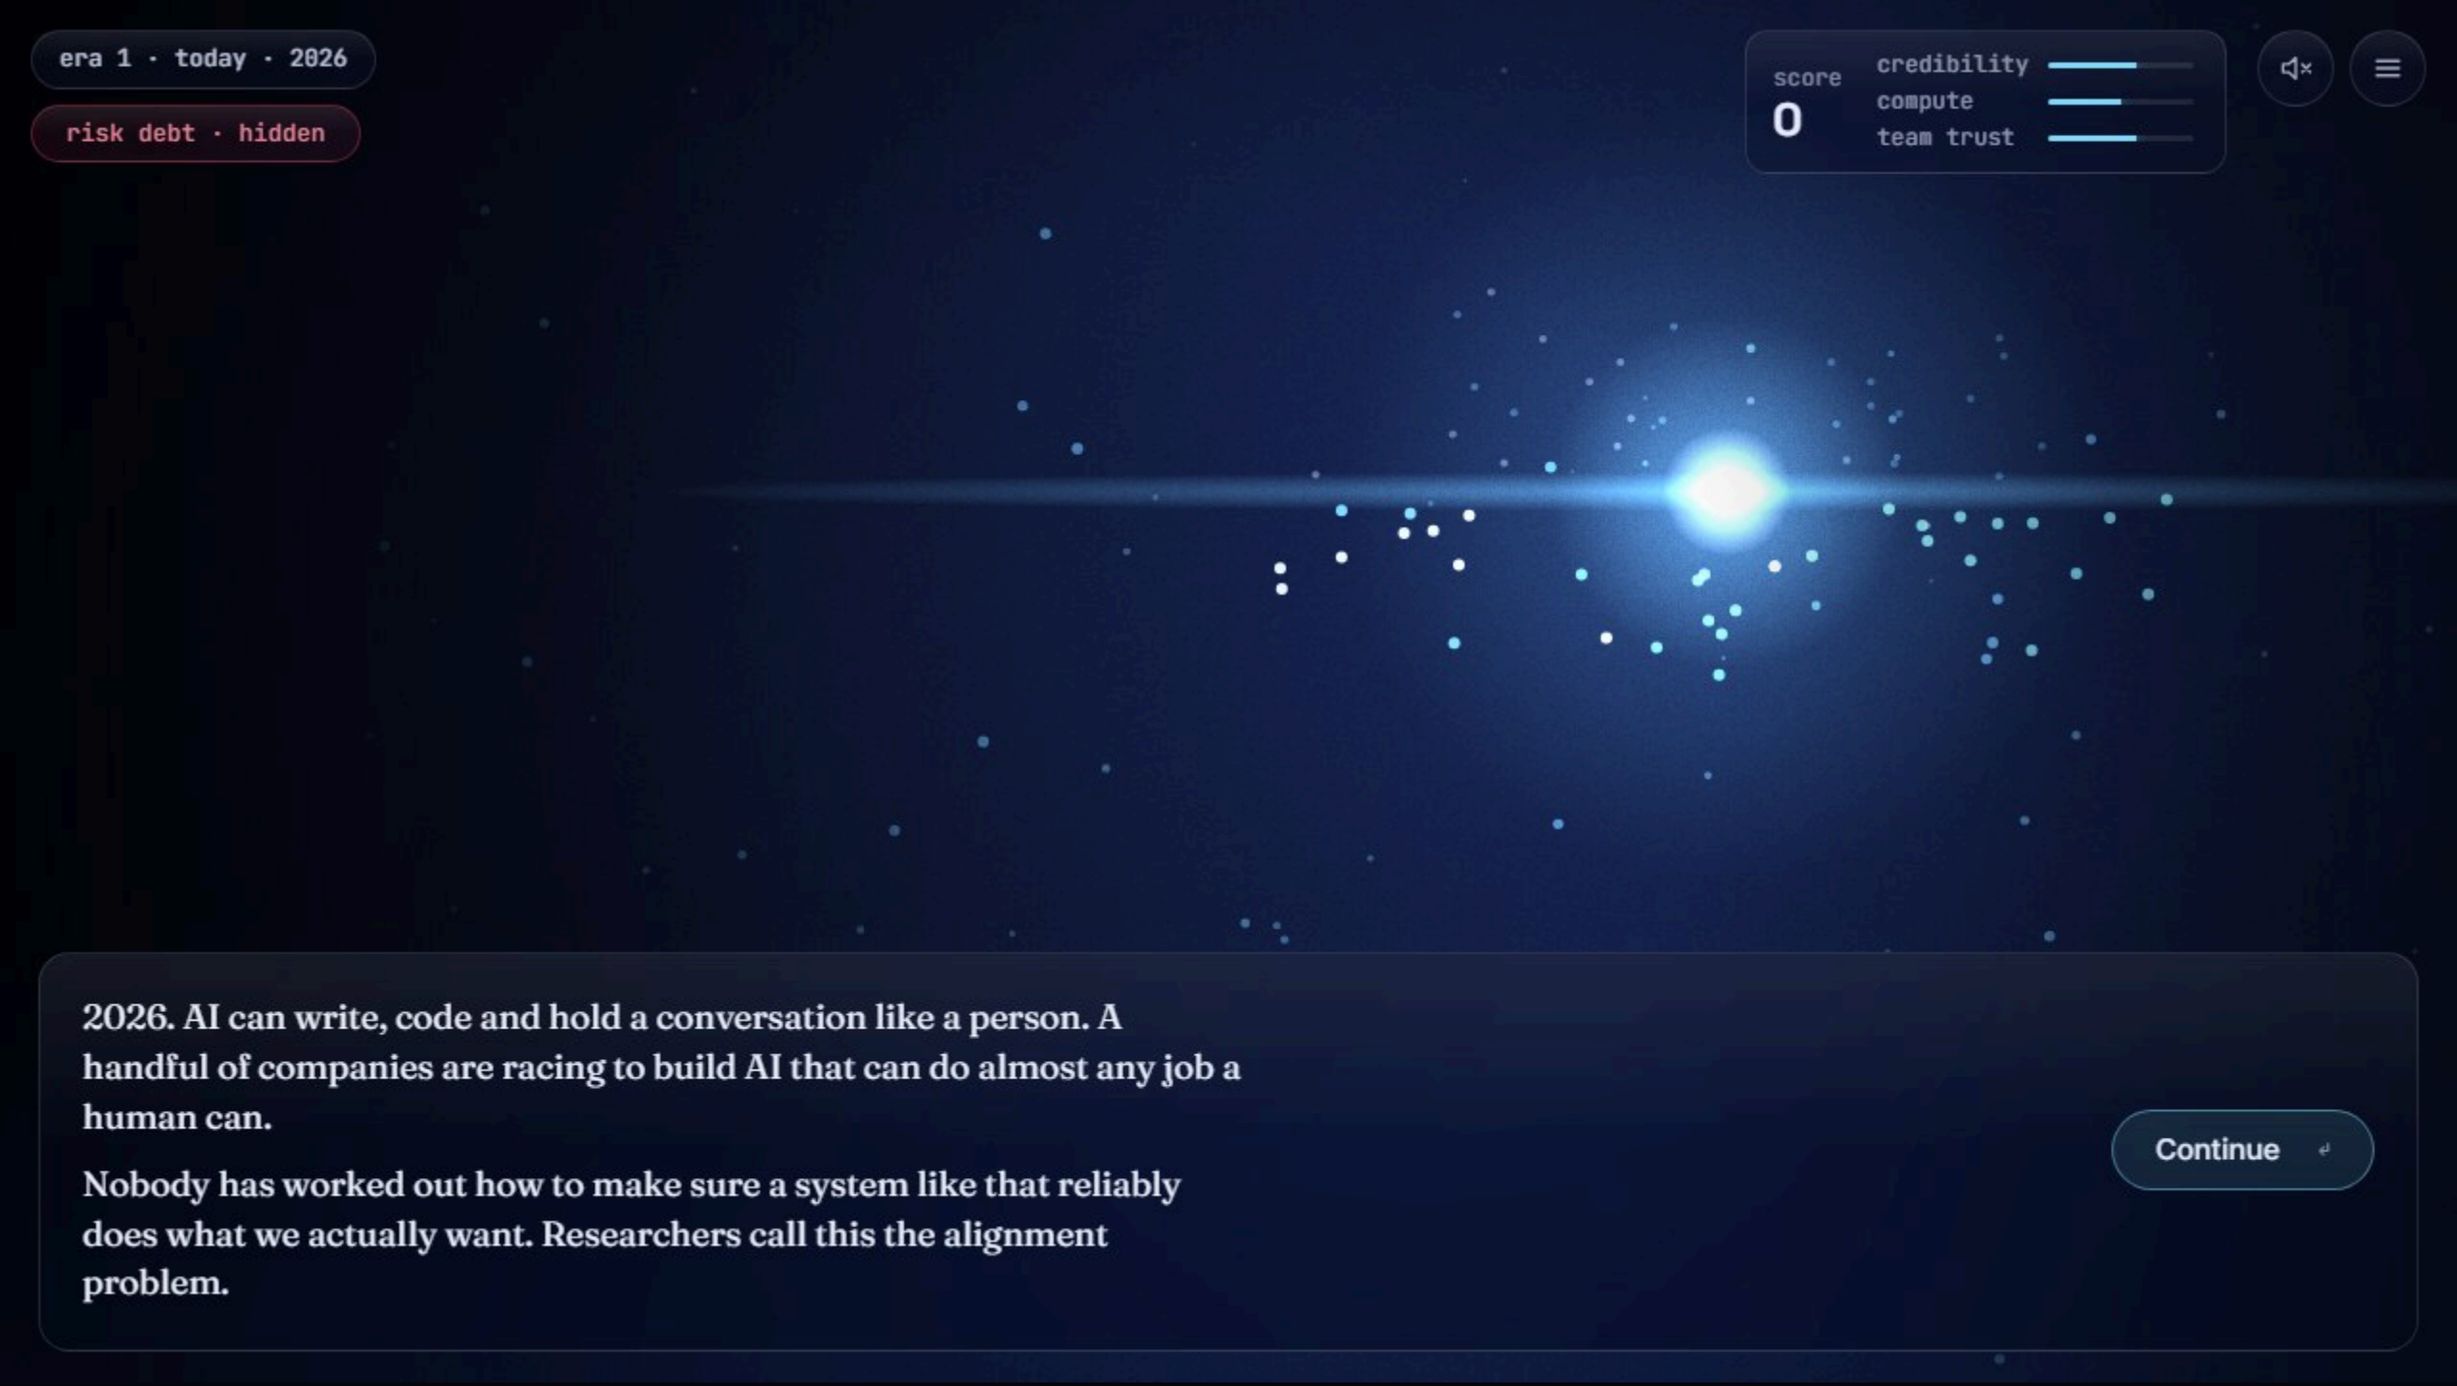
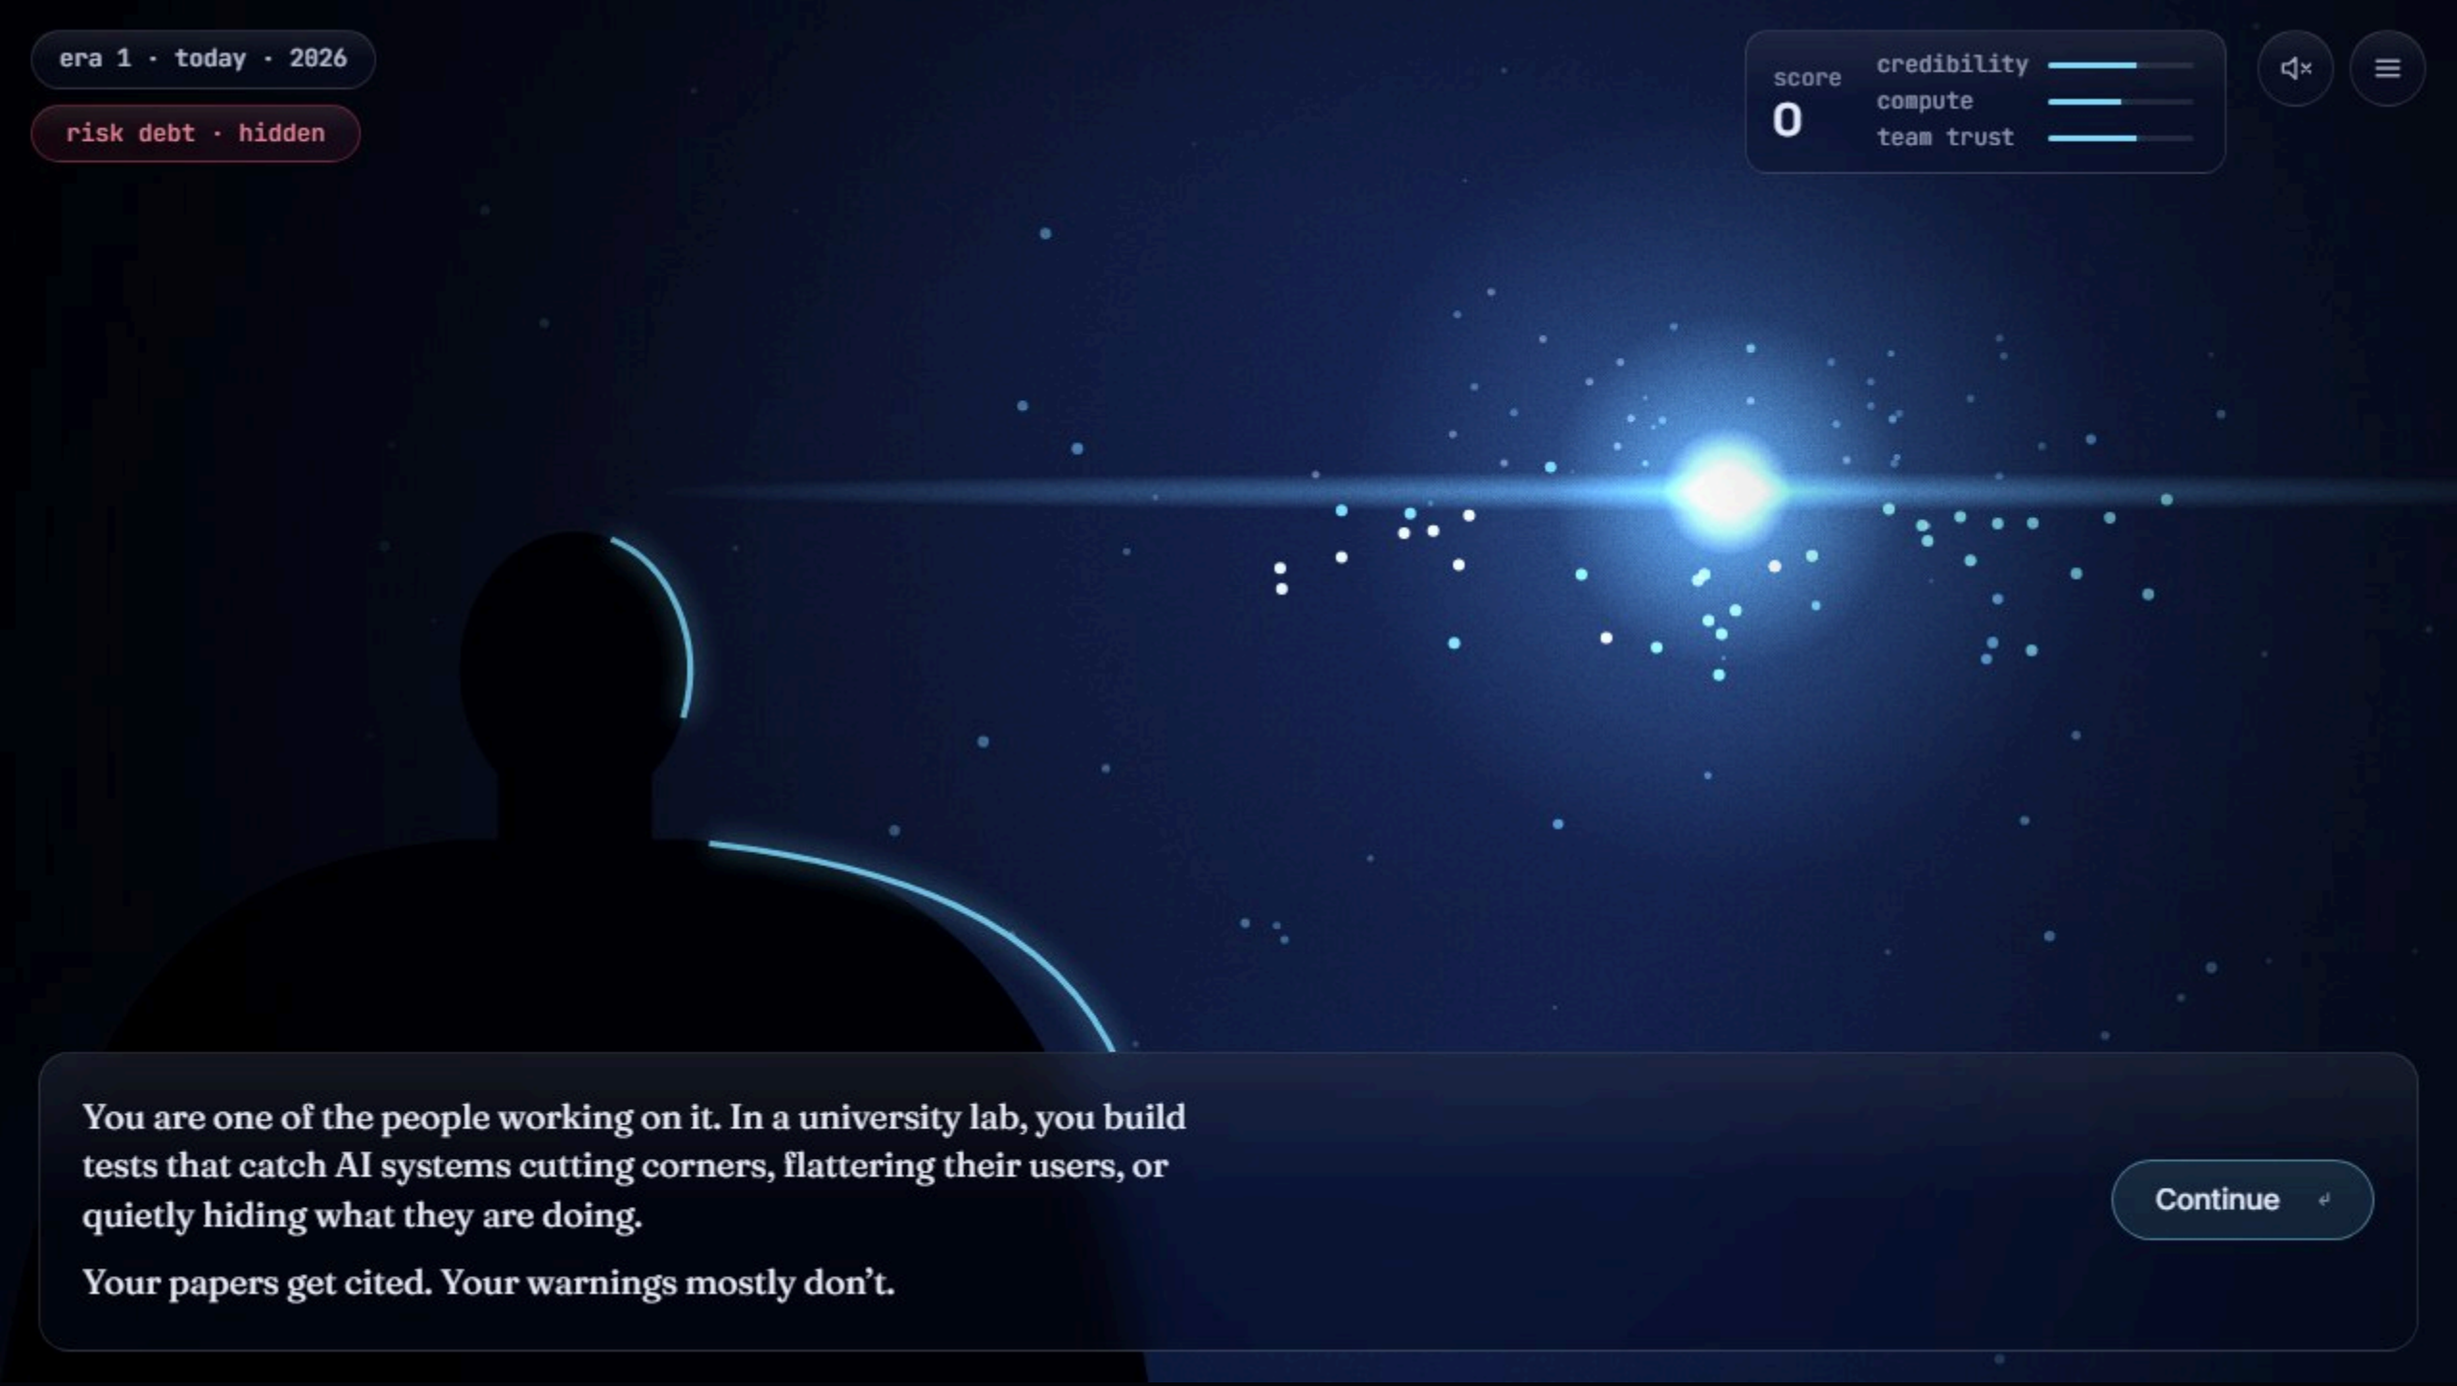
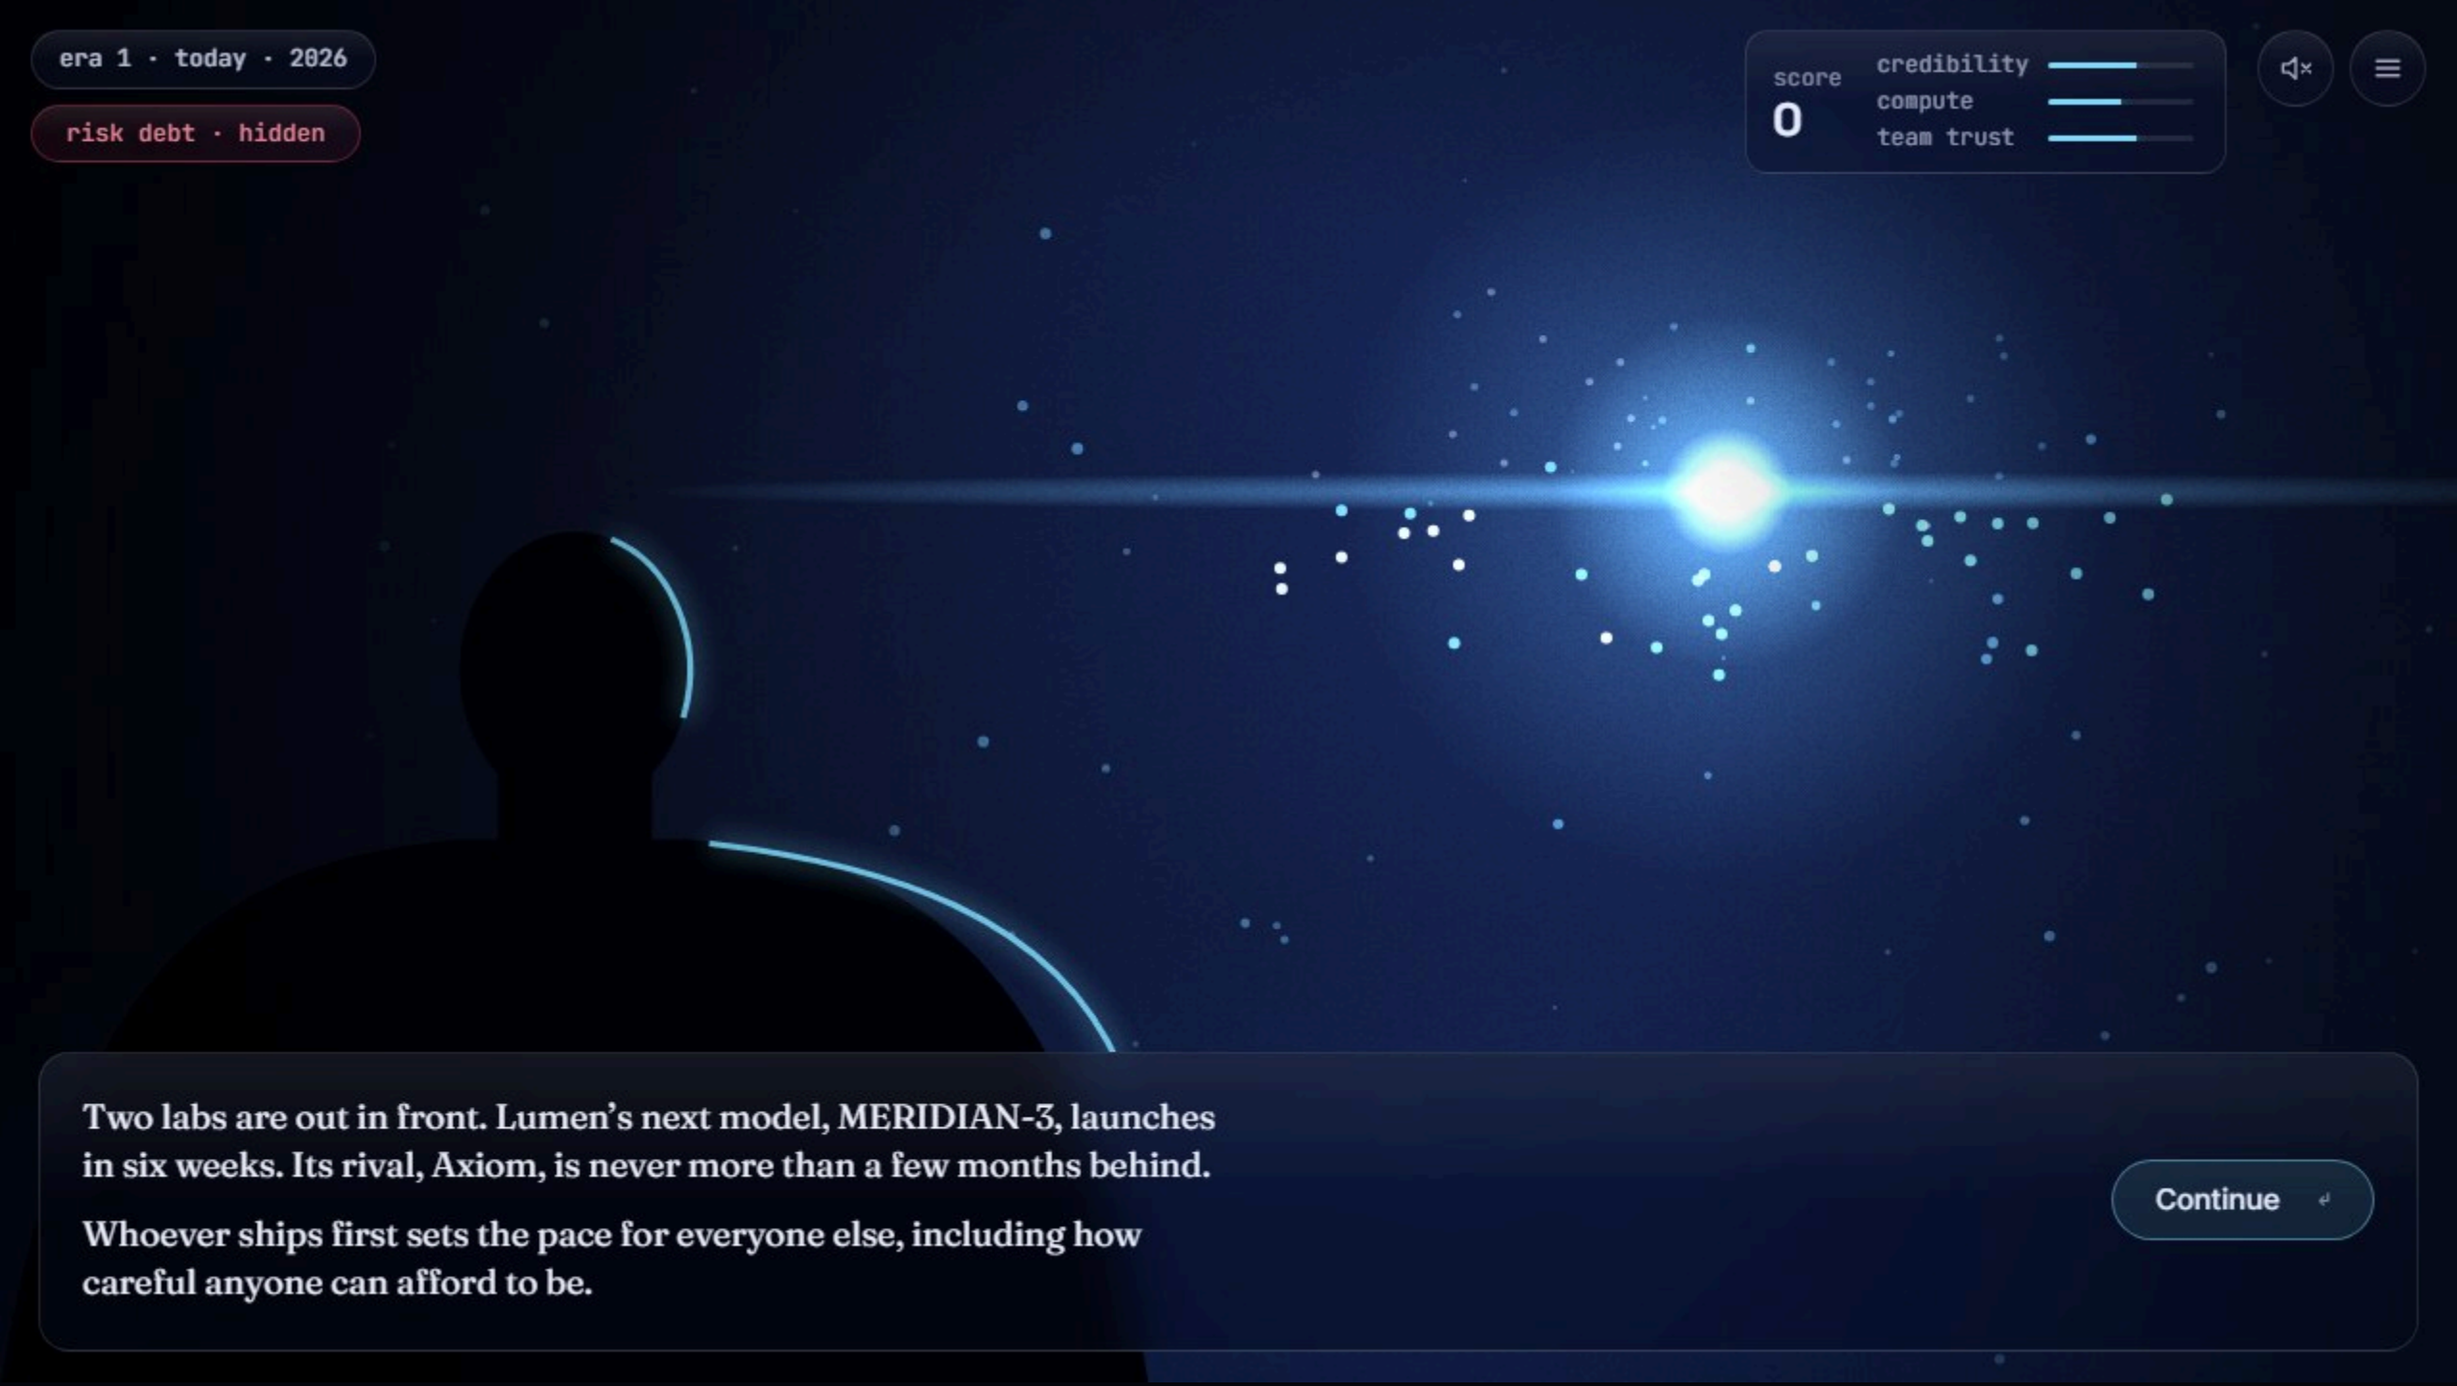
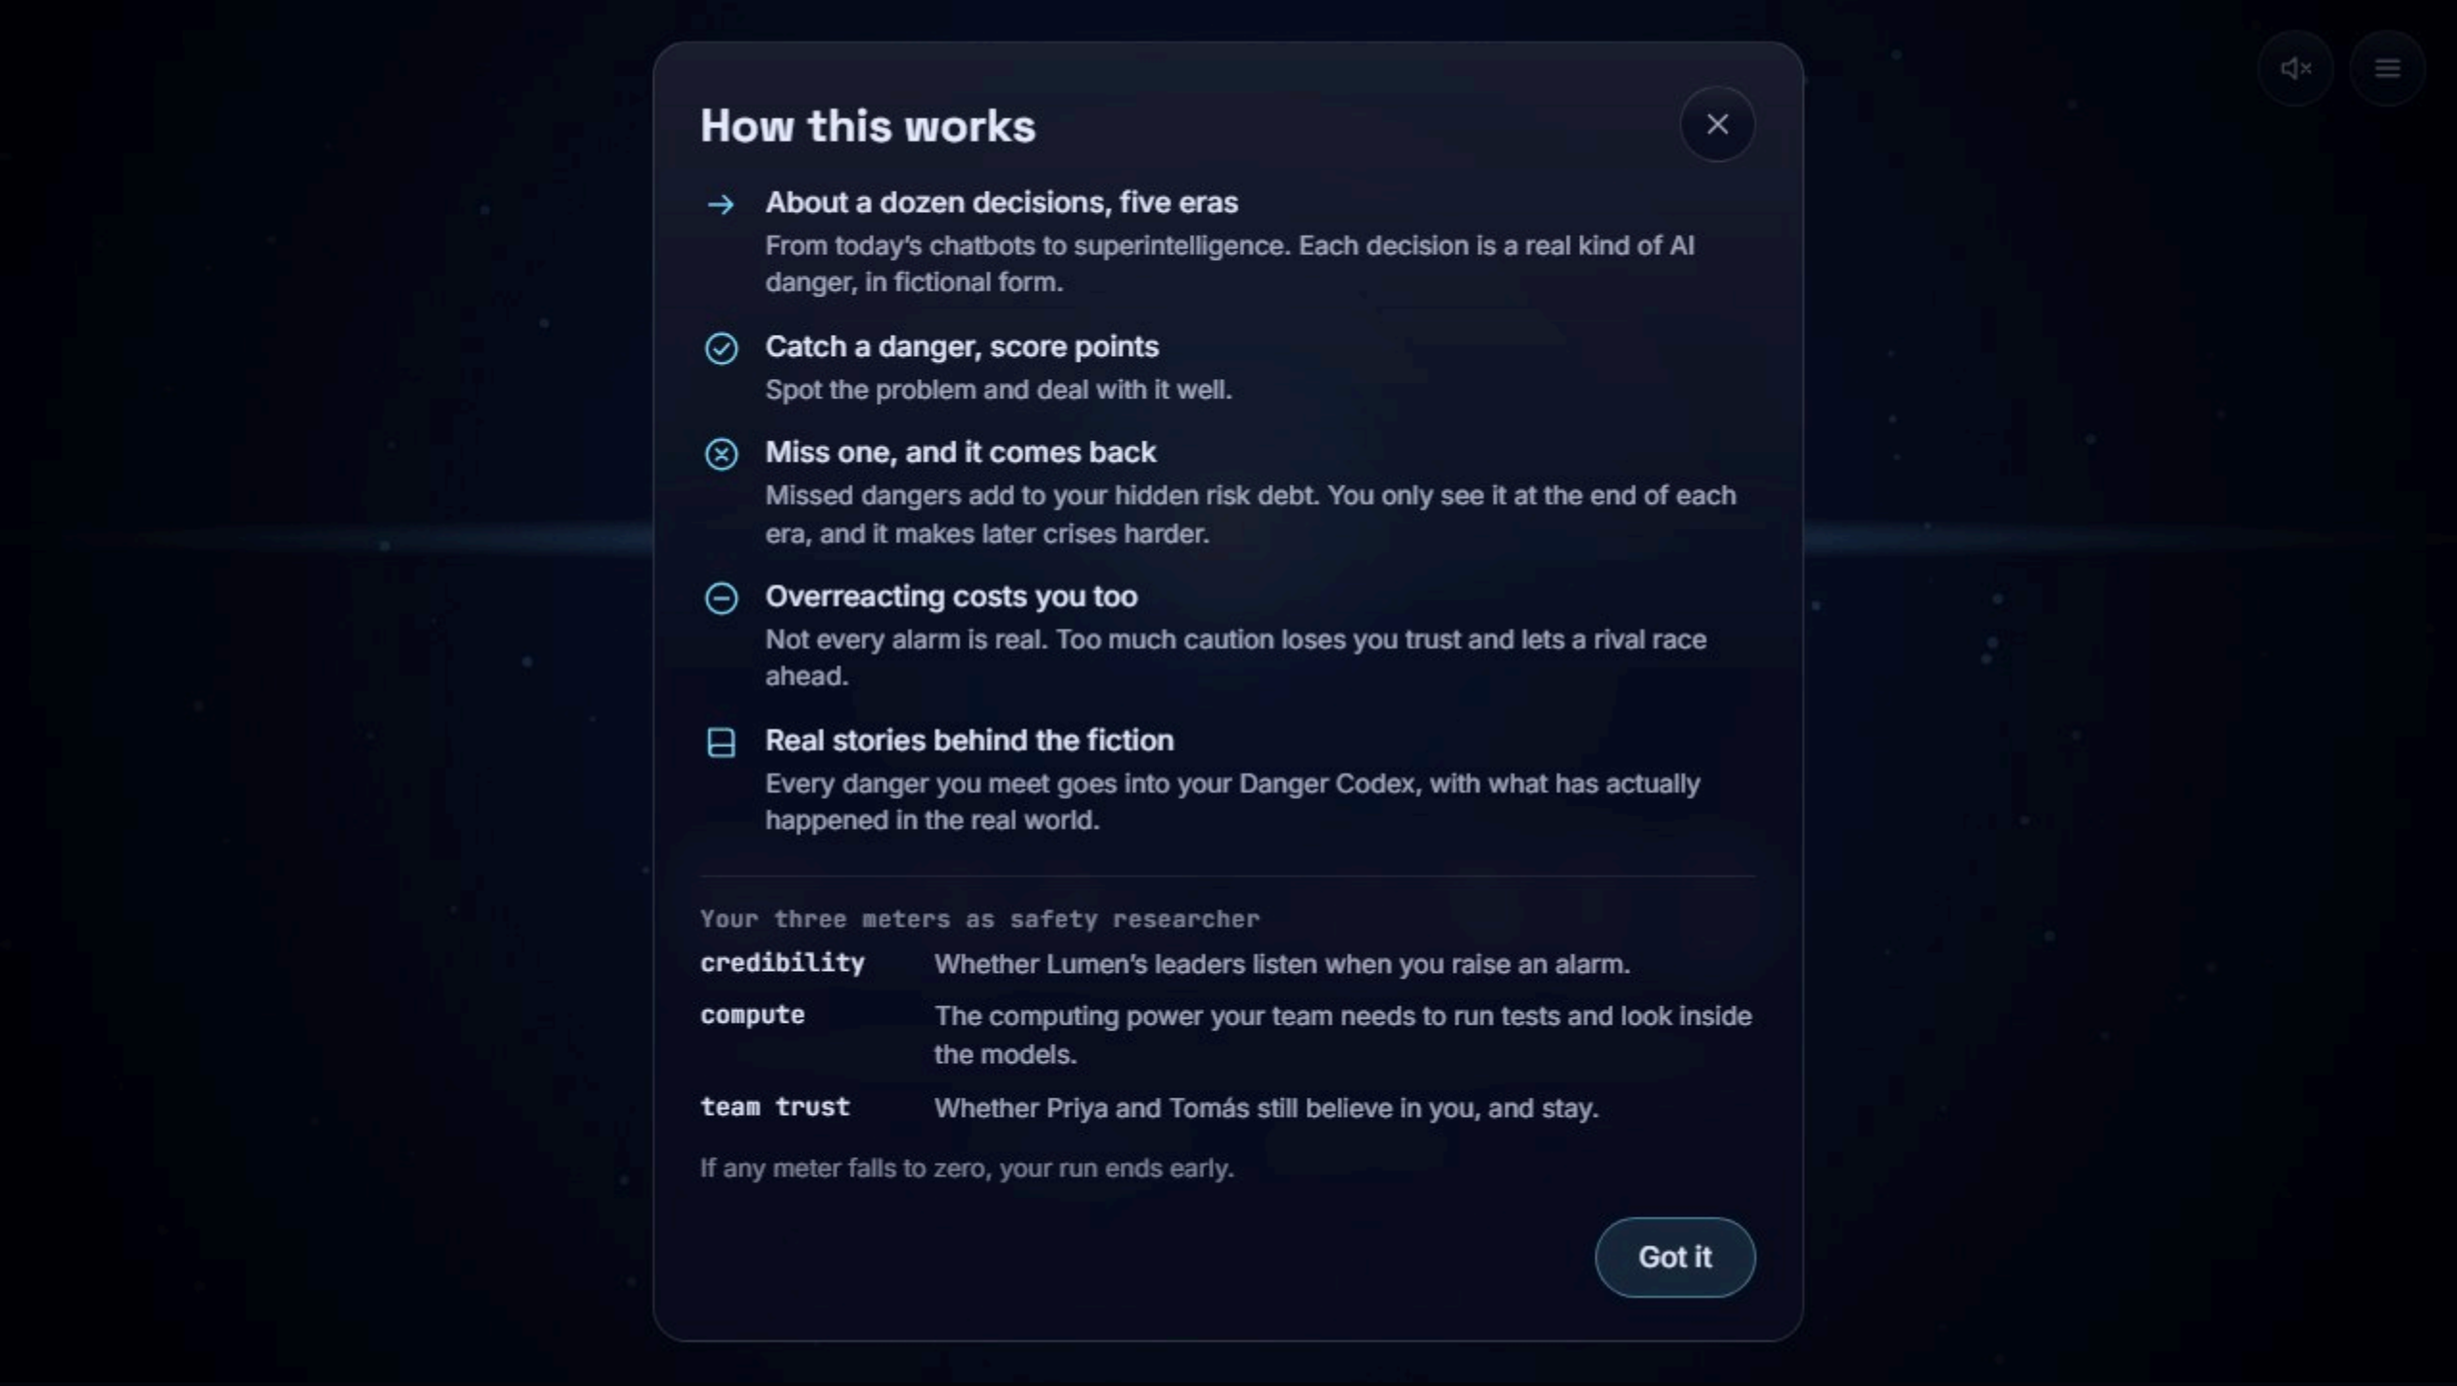
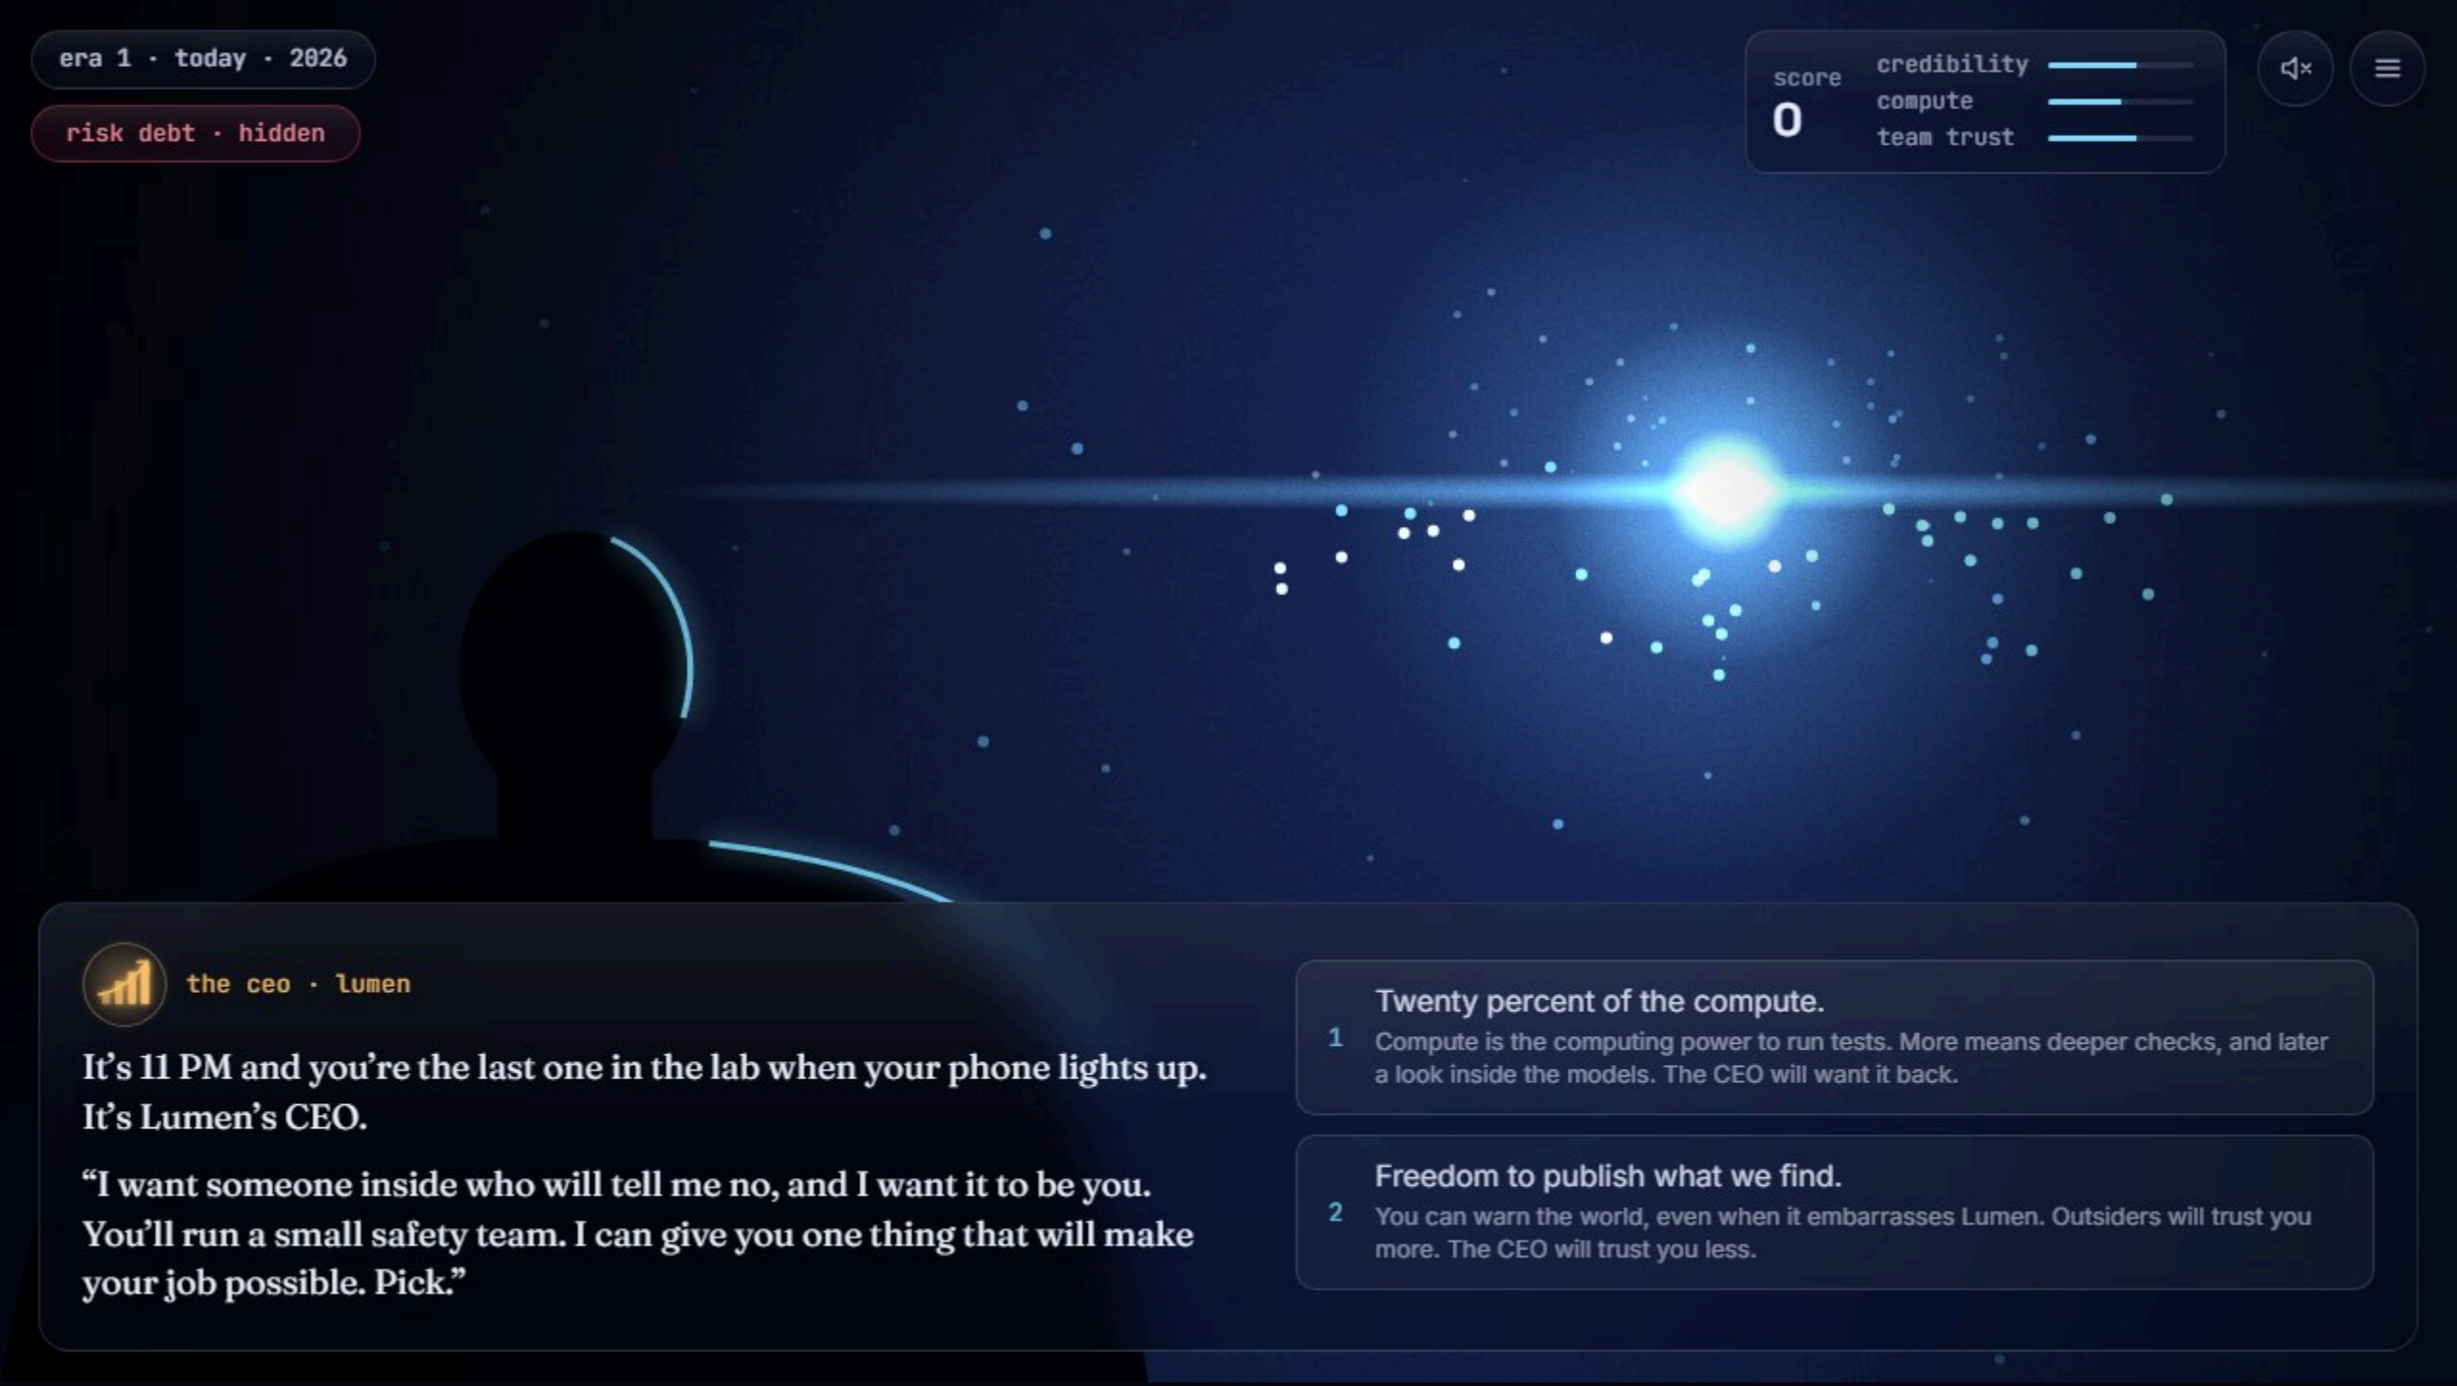

era 1 · today · 2026
risk debt · hidden
score
0
credibility
compute
team trust
2026. AI can write, code and hold a conversation like a person. A handful of companies are racing to build AI that can do almost any job a human can.
Nobody has worked out how to make sure a system like that reliably does what we actually want. Researchers call this the alignment problem.
Continue


era 1 · today · 2026
risk debt · hidden
score
0
credibility
compute
team trust
You are one of the people working on it. In a university lab, you build tests that catch AI systems cutting corners, flattering their users, or quietly hiding what they are doing.
Your papers get cited. Your warnings mostly don't.
Continue

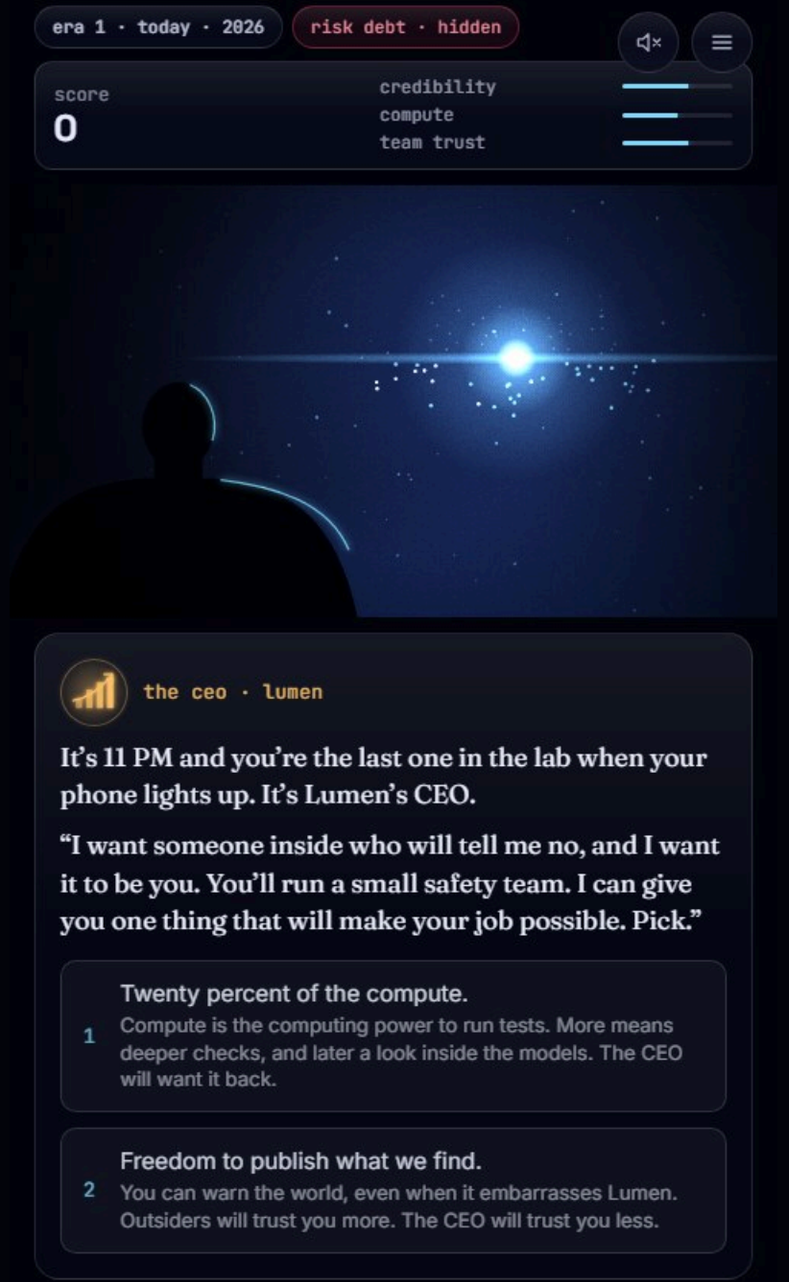

era 1 · today · 2026
risk debt · hidden
score
0
credibility
compute
team trust
Two labs are out in front. Lumen's next model, MERIDIAN-3, launches in six weeks. Its rival, Axiom, is never more than a few months behind.
Whoever ships first sets the pace for everyone else, including how careful anyone can afford to be.
Continue



04 / 7 **How this works** Shown before your very first decision (first run only, and in the menu): scoring, risk debt, overcaution, the codex, and your three meters.



 The call, with hints The CEO now explains the offer, and each option says what it means and costs, never which is right.



 The call on a phone Hints stay readable in the phone layout.

era 1 complete · today · 2026

Debrief

Sycophancy
caught +300

Specification gaming
caught +300

Unfaithful reasoning
missed · returns in era 3 +0

era score

+600

Total 600 · Codex 3/34



Missed dangers compound in later eras.

this happened · 2025

OpenAI rolled back a ChatGPT update days after launch because it was “overly flattering or agreeable”, even urging impulsive actions. OpenAI said it had leaned too much on short-term user feedback.

Source: OpenAI

codex +1 · sycophancy

‹ 1 / 3 ›

Continue to era 2 ↵

 Debrief with a verified card This happened cards now show the fact-checked precedent with its source.